Papers

1. Srivastava, Hinton, Krizhevsky, Sutskever, and Salakhutdinov. "Dropout: A Simple Way to Prevent Neural Networks from Overfitting." Journal of Machine Learning Research 15 (2014), pages 1929-1958.

2. Glorot, X., & Bengio, Y. Understanding the difficulty of training deep feedforward neural networks. AISTATS (2010), Vol. 9, pages 249-256.

3. Kingma and Ba. "ADAM: A method for stochastic optimization." arXiv preprint arXiv: 1412.6980 (2014).

4. "Simonyan and Zisserman. "Very deep convolutional networks for large-scale image recognition". ICLR, 2015.

5. Greff, Klaus, Rupesh Kumar Srivastava, Jan Koutník, Bas R. Steunebrink, and Jürgen Schmidhuber. "LSTM: A search space odyssey." arXiv preprint arXiv:1503.04069 (2015).

6. Chung, J., Gülçehre, C., Cho, K., & Bengio, Y. "Gated feedback recurrent neural networks." CoRR, 2015.

7. Jozefowicz, Zaremba, and Sutskever. "An Empirical Exploration of Recurrent Network Architectures." 2015.

8. Ioffe and Szegedy. "Batch Normalization: Accelerating Deep Network Training by Reducing Internal Covariate Shift". 2015.

9. Dean, Jeffrey, et al. "Large scale distributed deep networks." Advances in neural information processing systems. 2012.

10. Nemirovski, Juditsky, Lan, and Shapiro. "Robust Stochastic Approximation Approach to Stochastic Programming." SIAM J. Optim., 19(4), 1574–1609.

11. Auer, Peter, Nicolo Cesa-Bianchi, and Paul Fischer. "Finite-time analysis of the multiarmed bandit problem." Machine learning 47.2-3 (2002): 235-256.

12. Chapelle, Olivier, and Lihong Li. "An empirical evaluation of thompson sampling." Advances in neural information processing systems. 2011.

13. Watkins and Dayan. "Q-learning". 1992.

14. Mnih et al. "Playing Atari with Deep Reinforcement Learning". 2013.

15. Mnih, Volodymyr, et al. "Human-level control through deep reinforcement learning." Nature 518.7540 (2015): 529-533.

16. Osband et al. “Deep exploration via bootstrapped DQN”. Advances in neural information processing systems, 2016.

17. Osband and Van Roy. “Bootstrapped thompson sampling and deep exploration”. 2015.

18. Fortunato et al. “Noisy Networks for Exploration”. 2017

19. Hasselt, Guez, and Silver. “Deep Reinforcement Learning with Double Q-learning.” AAAI, 2016.

20. Wang et al. “Dueling network architectures for deep reinforcement learning.” 2015.

21. Lillicrap et al. “Continuous Control with Deep Reinforcement Learning.” ICLR, 2016.

22. Nair et al. “Massively parallel methods for deep reinforcement learning.” 2015.

23. Babaeizadeh et al. “GA3C: GPU-based A3C for deep reinforcement learning.” 2016.

24. Mnih et al. “Asynchronous methods for deep reinforcement learning”. ICML, 2016.

25. Schulman et al. “High-dimensional control using generalized advantage estimation.” 2015.