

IE 534/CS 598 Deep Learning

University of Illinois at Urbana-Champaign

Fall 2018

Lecture 1

“Most of all, there is a shortage of [deep learning] talent, and the big companies are trying to land as much of it as they can. Solving tough A.I. problems is not like building the flavor-of-the-month smartphone app. In the entire world, fewer than 10,000 people have the skills necessary to tackle serious artificial intelligence research [...]”

– New York Times, October 2017.

Topics:

- Fully-connected networks
- Convolution networks
- Residual networks
- Recurrent networks (e.g., LSTM)
- Deep reinforcement learning
- Generative adversarial networks
- Optimization and training methods

PyTorch is a software library for training deep learning models.

- Seamlessly integrated with Python.
- **Define-by-run** framework allows for dynamic training of models.
- Capable of distributing training across multiple machines.

Computational resources:

- 50,000 GPU hours
- Unique opportunity to implement large-scale deep learning models!
- Training of deep learning models can be highly parallelized on GPUs (frequently $10\times$ faster than CPUs).

Homeworks

- 1 Implement a neural network from scratch in Python for the MNIST dataset (no PyTorch).
- 2 Implement a convolution neural network from scratch in Python for the MNIST dataset (no PyTorch).
- 3 Deep convolution network for MNIST and CIFAR10.
- 4 Residual neural network for CIFAR100.
- 5 Image ranking with convolution networks.
- 6 Natural Language Processing I
- 7 Natural Language Processing II
- 8 Generative adversarial networks (GANs)
- 9 Video recognition.

Extensive Code and Notes will be provided, including:

PyTorch (and TensorFlow) Code

Documentation/code on using PyTorch on Blue Waters

Course notes

Grading:

35% Homeworks

35% Midterm (**November 6, 8:00 AM**)

30% Final Project

Course website: <https://courses.engr.illinois.edu/ie534/fa2018/>

Submit homeworks via Compass (due on Fridays at **5:00 PM**).

Start homeworks early!

4 office hours per week.

- Machine learning estimates a statistical model for the relationship between an input X and an output Y .
- Formally, suppose there is data $(X, Y) \in \mathbb{R}^d \times \mathcal{Y}$ and a statistical model $f(x; \theta) : \mathbb{R}^d \rightarrow \mathbb{R}^K$.
- $\theta \in \Theta$ are the parameters in the model and must be estimated.
- We wish to find a model $f(x; \theta)$ such that $f(X; \theta)$ is “an accurate prediction” for Y .

$$\mathcal{L}(\theta) = \mathbb{E}_{(X,Y)} [\rho(f(X; \theta), Y)]. \quad (1)$$

- $\rho(z, y)$ measures the distance between the model prediction z and y .
- This distance is then averaged over the distribution $\mathbb{P}_{(X,Y)}$ of the data (X, Y) .

The best model, within the class of models $\{f(x; \theta')\}_{\theta' \in \Theta}$, is the model $f(x; \theta)$ where θ satisfies

$$\theta = \arg \min_{\theta' \in \Theta} \mathcal{L}(\theta'). \quad (2)$$

- Typically, the distribution $\mathbb{P}_{(X,Y)}$ is unknown.
- Instead, i.i.d. data samples $(x^n, y^n)_{n=1}^N$ are available from the distribution $\mathbb{P}_{(X,Y)}$.
- Then, our objective function becomes

$$\mathcal{L}^N(\theta) = \frac{1}{N} \sum_{n=1}^N \rho(f(x^n; \theta), y^n). \quad (3)$$

- As the number of data samples $N \rightarrow \infty$, $\mathcal{L}^N(\theta) \rightarrow \mathcal{L}(\theta)$.

The best model is

$$\theta = \arg \min_{\theta' \in \Theta} \mathcal{L}^N(\theta'). \quad (4)$$

- For complicated models such as neural networks, these minimization problems cannot be exactly calculated.
- Instead, numerical methods are used to minimize the objective functions.
- Convexity versus Non-Convexity
- In the non-convex case, numerical methods are only guaranteed to converge to a point which satisfies certain optimization properties.
- **Stochastic gradient descent** is the method of choice for training deep learning models.

Example

- Consider a logistic regression model for classification where $\mathcal{Y} = \{0, 1, \dots, K - 1\}$ and $\Theta = \mathbb{R}^{K \times d}$.
- Given an input $x \in \mathbb{R}^d$, the model $f(x; \theta)$ produces a probability of each possible outcome in \mathcal{Y} :

$$\begin{aligned} f(x; \theta) &= F_{\text{softmax}}(\theta x), \\ F_{\text{softmax}}(z) &= \frac{1}{\sum_{k=0}^{K-1} e^{z_k}} \left(e^{z_0}, e^{z_1}, \dots, e^{z_{K-1}} \right). \end{aligned} \quad (5)$$

- $F_{\text{softmax}}(z)$ takes a K -dimensional input and produces a probability distribution on \mathcal{Y} .
- The function $F_{\text{softmax}}(z) : \mathbb{R}^K \rightarrow \mathcal{P}(\mathcal{Y})$ is called the “softmax function” and is frequently used in deep learning.

The objective function is the negative log-likelihood (commonly referred to in machine learning as the “cross-entropy error”):

$$\begin{aligned}\mathcal{L}(\theta) &= \mathbb{E}_{(X,Y)}[\rho(f(X;\theta), Y)], \\ \rho(z, y) &= - \sum_{k=0}^{K-1} \mathbf{1}_{y=k} \log z_k,\end{aligned}\tag{6}$$

where z_k is the k -th element of the vector z and $\mathbf{1}_{y=k}$ is the indicator function

$$\mathbf{1}_{y=k} = \begin{cases} 1 & y = k \\ 0 & y \neq k \end{cases}$$

Gradient Descent (GD):

$$\theta^{(\ell+1)} = \theta^{(\ell)} - \alpha^{(\ell)} \nabla_{\theta} \mathcal{L}(\theta^{(\ell)}). \quad (7)$$

- Gradient descent repeatedly takes steps in the direction of *steepest descent*.
- The magnitude of these steps is governed by the “learning rate” $\alpha^{(\ell)}$, which is a positive scalar which may depend upon the iteration number ℓ .

We can show that if the learning rate $\alpha^{(\ell)}$ is sufficiently small, the ℓ -th step of the gradient descent algorithm (7) is guaranteed to decrease the objective function.

Using a Taylor expansion,

$$\begin{aligned}\mathcal{L}(\theta^{(\ell+1)}) - \mathcal{L}(\theta^{(\ell)}) &= \nabla_{\theta} \mathcal{L}(\theta^{(\ell)}) (\theta^{(\ell+1)} - \theta^{(\ell)}) \\ &+ \frac{1}{2} (\theta^{(\ell+1)} - \theta^{(\ell)})^{\top} \nabla_{\theta\theta} \mathcal{L}(\bar{\theta}) (\theta^{(\ell+1)} - \theta^{(\ell)}).\end{aligned}\quad (8)$$

Substitute for $\theta^{(\ell+1)} - \theta^{(\ell)}$ using the gradient descent update equation:

$$\begin{aligned}\mathcal{L}(\theta^{(\ell+1)}) - \mathcal{L}(\theta^{(\ell)}) &= -\alpha^{(\ell)} \left(\nabla_{\theta} \mathcal{L}(\theta^{(\ell)}) \right)^{\top} \nabla_{\theta} \mathcal{L}(\theta^{(\ell)}) \\ &+ \frac{1}{2} \left(\alpha^{(\ell)} \right)^2 \nabla_{\theta} \mathcal{L}(\theta^{(\ell)})^{\top} \nabla_{\theta\theta} \mathcal{L}(\bar{\theta}) \nabla_{\theta} \mathcal{L}(\theta^{(\ell)}),\end{aligned}\quad (9)$$

- It is also clear that if $\alpha^{(\ell)}$ is too large, the objective function may *increase* due to the second-order term.
- In practice, a careful choice of the learning rate is very important.
- The gradient descent algorithm uses only the first derivative $\nabla_{\theta}\mathcal{L}(\theta)$ to update the parameter θ .
- If it takes too large of a step, the first derivative no longer accurately describes the change in the objective function.

Gradient descent requires computing the gradient $\nabla_{\theta}\mathcal{L}(\theta^{(\ell)})$, which can be computationally costly since it involves an integral over (x, y) :

$$\begin{aligned}\nabla_{\theta}\mathcal{L}(\theta^{(\ell)}) &= \nabla_{\theta}\mathbb{E}_{(X,Y)}[\rho(f(X;\theta^{(\ell)}), Y)] \\ &= \mathbb{E}_{(X,Y)}[\nabla_{\theta}\rho(f(X;\theta^{(\ell)}), Y)].\end{aligned}\tag{10}$$

Stochastic gradient descent (SGD) is a computationally efficient scheme for minimizing $\mathcal{L}(\theta)$.

It follows a *noisy* (but unbiased) descent direction:

$$\theta^{(\ell+1)} = \theta^{(\ell)} - \alpha^{(\ell)} \nabla_{\theta} \rho(f(x^{(\ell)}; \theta^{(\ell)}), y^{(\ell)}), \quad (11)$$

where $(x^{(\ell)}, y^{(\ell)})$ are i.i.d. samples from the distribution $\mathbb{P}_{(X, Y)}$.

The *average* descent direction in (11) equals the GD algorithm's descent direction since

$$\begin{aligned} & \mathbb{E} \left[\nabla_{\theta} \rho(f(x^{(\ell)}; \theta^{(\ell)}), y^{(\ell)}) \middle| \theta^{(\ell)} \right] \\ &= \mathbb{E} \left[\nabla_{\theta} \rho(f(X; \theta^{(\ell)}), Y) \middle| \theta^{(\ell)} \right] \\ &= \nabla_{\theta} \mathcal{L}(\theta^{(\ell)}). \end{aligned} \quad (12)$$

The distribution $\mathbb{P}_{(X,Y)}$ is usually unknown.

Instead, data samples $(x_n, y_n)_{n=1}^N$ are available from the distribution $\mathbb{P}_{(X,Y)}$.

Then, objective function can be approximated as

$$\mathcal{L}(\theta) = \frac{1}{N} \sum_{n=1}^N \rho(f(x^n; \theta), y^n). \quad (13)$$

The gradient descent algorithm for (13) is

$$\theta^{(\ell+1)} = \theta^{(\ell)} - \alpha^{(\ell)} \frac{1}{N} \sum_{n=1}^N \nabla_{\theta} \rho(f(x^n; \theta^{(\ell)}), y^n). \quad (14)$$

The stochastic gradient descent algorithm is:

- Randomly initialize the parameter $\theta^{(0)}$.
- For $\ell = 0, 1, \dots, L$:
 - Select a data sample $(x^{(\ell)}, y^{(\ell)})$ at random from the dataset $(x_n, y_n)_{n=1}^N$.
 - Calculate the gradient for the loss from the data sample $(x^{(\ell)}, y^{(\ell)})$:

$$G^{(\ell)} = \nabla_{\theta} \rho(f(x^{(\ell)}; \theta^{(\ell)}), y^{(\ell)}) \quad (15)$$

- Update the parameters:

$$\theta^{(\ell+1)} = \theta^{(\ell)} - \alpha^{(\ell)} G^{(\ell)}, \quad (16)$$

where $\alpha^{(\ell)}$ is the learning rate.

- **In practice, stochastic gradient descent typically converges much more rapidly than gradient descent!**
- GD converges slowly since in order to take a single step, it must calculate the gradients for every data sample in the dataset.
- In contrast, SGD can rapidly take many steps since each step only requires calculating the gradient for a single data sample.
- **SGD is especially advantageous when the size of the dataset N is large.**

The learning rate must satisfy the following conditions in order for SGD to converge:

$$\begin{aligned}\sum_{\ell=0}^{\infty} \alpha^{(\ell)} &= \infty, \\ \sum_{\ell=0}^{\infty} (\alpha^{(\ell)})^2 &< \infty.\end{aligned}\tag{17}$$

A learning rate which satisfies these conditions is

$$\alpha^{(\ell)} = \frac{C_0}{C_1 + \ell}.\tag{18}$$

It is often sufficient in practice to simply use a piecewise learning rate schedule for $\ell = 0, 1, \dots, K_4$ such as

$$\alpha^{(\ell)} = \begin{cases} C & \ell \leq K_1 \\ C \times 10^{-1} & K_1 < \ell \leq K_2 \\ C \times 10^{-2} & K_2 < \ell \leq K_3 \\ C \times 10^{-3} & K_3 < \ell \leq K_4 \end{cases}$$

If the learning rate is too small, convergence may be very slow.

If the learning rate is too large, the algorithm may oscillate and make no progress.

Theorem

Suppose that $\nabla_{\theta}\mathcal{L}(\theta)$ is globally Lipschitz and bounded. Furthermore, assume that the condition (17) holds and $\mathcal{L}(\theta)$ is bounded. Then,

$$\mathbb{P}\left[\lim_{\ell \rightarrow \infty} \nabla_{\theta}\mathcal{L}(\theta^{(\ell)}) = 0\right] = 1.$$

- Neural networks are not globally Lipschitz.
- Neural networks are not bounded.
- Neural networks are non-convex: SGD may converge to a local minimum and not a global minimum!

The **mini-batch stochastic gradient descent algorithm** is:

- Randomly initialize the parameter $\theta^{(0)}$.
- For $\ell = 0, 1, \dots, L$:
 - Select M data samples $(x^{(\ell,m)}, y^{(\ell,m)})_{m=1}^M$ at random from the dataset $(x_n, y_n)_{n=1}^N$, where $M \ll N$.
 - Calculate the gradient for the loss from the data samples:

$$G^{(\ell)} = \frac{1}{M} \sum_{m=1}^M \nabla_{\theta} \rho(f(x^{(\ell,m)}; \theta^{(\ell)}), y^{(\ell,m)}) \quad (19)$$

- Update the parameters:

$$\theta^{(\ell+1)} = \theta^{(\ell)} - \alpha^{(\ell)} G^{(\ell)}, \quad (20)$$

where $\alpha^{(\ell)}$ is the learning rate.

- The mini-batch update $G^{(\ell)}$ is clearly still an unbiased estimate for the gradient $\nabla_{\theta}\mathcal{L}(\theta^{(\ell)})$.
- It is less noisy than the stochastic gradient descent update with a single sample, i.e.

$$\begin{aligned}\text{Var}\left[G^{(\ell)}\middle|\theta^{(\ell)}\right] &= \text{Var}\left[\frac{1}{M}\sum_{m=1}^M\nabla_{\theta}\rho(f(x^{(\ell,m)};\theta^{(\ell)}),y^{(\ell,m)})\middle|\theta^{(\ell)}\right] \\ &= \frac{1}{M}\text{Var}\left[\nabla_{\theta}\rho(f(x^{(\ell)};\theta^{(\ell)}),y^{(\ell)})\middle|\theta^{(\ell)}\right].\end{aligned}\quad (21)$$

- The conditional variance of a mini-batch update is smaller by a factor of $\frac{1}{M}$ than stochastic gradient descent with a single sample, where M is the mini-batch size.